

Covering the Databricks Lakehouse platform — releases, capabilities, and engineering decisions behind them.

Technical Features

NOW GENERALLY AVAILABLE

Catalog Commits for Unity Catalog

Catalog Commits makes Delta and Unity Catalog **co-evolve as a unified system**.

External engines can now write to UC-managed Delta tables with full guarantees. Multi-table **ACID transactions** are consistent across all engines, with governance and access control enforced automatically.

BUSINESS VALUE

Eliminates data silos between tools and compute engines. Any external system can now read from and write to governed Lakehouse data with consistent access control. This is foundational for multi-cloud and multi-engine architectures where enforcing a single governance layer across all producers has historically required custom workarounds.

→ [Databricks docs](#)

BETA · PIPELINE EFFICIENCY

SDP REPLACE WHERE — Selective Recomputation

SDP now supports **REPLACE WHERE** — rewriting only the rows matching a predicate (e.g. the last 30 days) without touching the rest of the table.

Combined with Pipeline Parameters, teams can now **parameterise exactly which date range or partition to refresh** — making a single pipeline definition reusable across both incremental and backfill scenarios.

ENGINEERING SPOTLIGHT

Why REPLACE WHERE changes how you think about pipelines ?

REPLACE WHERE rewrites only matching rows using standard SQL or Python — no full rebuilds, no complex partition scripts.

BUSINESS VALUE

Dramatically reduces pipeline compute costs. Backfills, dimension changes, and late-arriving data no longer require full-table rebuilds. Pipeline costs become predictable and proportional to the scope of change, not to table size.

→ [Databricks docs](#)

BETA · LAKEFLOW

Pipeline Parameters in Lakeflow Declarative Pipelines

Lakeflow now supports **configurable runtime parameters in Lakeflow Declarative Pipelines**. Teams can define key-value pairs — such as `environment: dev` or `catalog: bronze` — that are resolved at runtime, eliminating hardcoded values that require separate pipeline definitions per environment or workload.

Changes propagate from one source of truth, and promotion no longer means maintaining separate pipeline definitions per environment.

BUSINESS VALUE

One pipeline definition per workload, across all environments. Reduces maintenance overhead and accelerates promotion cycles from development to production. Pipelines become genuinely reusable assets rather than environment-specific configurations.

→ [Databricks docs](#)

PRODUCTIVITY · NOTEBOOKS

%uv pip — Faster Package Management in Notebooks

The new `%uv pip` command brings the `uv` package manager to Databricks notebooks, delivering significantly faster package installs and dependency resolution than the existing `%pip` approach, while using the same virtual environment — making it a drop-in replacement with no workflow changes required. Currently available on Serverless Environment 5 only. Best suited to interactive iteration workflows; not yet recommended for reproducible production environments.

→ [Databricks docs](#)

LAKEFLOW JOBS · GENERALLY AVAILABLE

Disabled Tasks in Lakeflow Jobs

A `disabled: true` property can now be set on any task within a Lakeflow job definition, either via the UI or YAML/DABs configuration. Disabled tasks — and all downstream dependents — are treated as skipped rather than failed, allowing the job to complete successfully without touching the broader DAG.

This addresses a persistent pain point in pipeline debugging, temporary backfills, and partial testing workflows where the only previous option was to restructure or fork the job definition.

→ [Databricks docs](#)

INCREMENTAL INGESTION · NO CDF REQUIRED

Watermark-Based Query Capture

Change Data Feed is not the only path to incremental ingestion. Lakeflow now supports **query-based watermark capture** as a CDF-free pattern: pull rows where a timestamp column like `last_modified_ts` or `updated_at` — exceeds the stored watermark. Soft deletes use `is_deleted = TRUE`, and the watermark advances on success. Works well with SCD Type 1 targets where CDF is unavailable.

→ [Databricks docs](#)

SPARK 4.X · DEVELOPER EXPERIENCE

Python Data Source API — Custom Connectors in ~20 Lines

Spark 4's **Python Data Source API** collapses custom connector development from a 2–6 week Scala or Java project to a single Python class of approximately 20 lines plus one registration call. Now Generally Available on Databricks, connectors are registered once and consumed via the standard `spark.read.format(...)` interface — indistinguishable from native sources.

[→ Databricks docs](#)

SERVERLESS · INFRASTRUCTURE

Serverless Workspaces — Fully Managed

A **Serverless Workspace** comes pre-configured with serverless compute and default storage, with Databricks managing all underlying cloud infrastructure — VPCs, compute instances, and storage wiring. It delivers a **SaaS-style Databricks experience** — Databricks handles all infrastructure so teams can start immediately.

[→ Databricks docs](#)

Covering the Databricks Lakehouse platform — releases, capabilities, and engineering decisions behind them.

In-General / Business-Focused Features

MAJOR RELEASE

Databricks DQX v0.14.0

DQX v0.14.0 introduces **DQX Studio** — a browser app that lets analysts author and run data quality rules without writing code.

It also adds **ML-based row-level anomaly detection**, plus **versioned fingerprints** for auditability, template variables for reuse across environments, and three new built-in checks:

`has_no_aggr_outliers`, `are_polygons_mutually_disjoint`, and `foreign_key_null_safe`. AI Agent Skills for Genie Code, Claude Code, and Cursor round out the integration story.

BUSINESS VALUE

Reduces dependency on engineers for data quality. Business analysts can now define, test, and deploy quality rules through a UI. ML-driven anomaly detection replaces fragile threshold configs. Full rule versioning and fingerprinting satisfies audit requirements in regulated industries — without additional tooling.

[→ Databricks docs](#)

NEW INTEGRATIONS

Excel Add-in & Google Sheets Connector

The Excel Add-in and Google Sheets Connector bring **live, governed Lakehouse data directly into the spreadsheet tools that business teams already use daily** — no SQL or CSV exports required.

Unity Catalog permissions are enforced automatically — users see only what they're entitled to, with no extra access management overhead.

BUSINESS VALUE

Unlocks data access for finance, operations, and sales teams. Eliminates stale CSV exports, duplicated metrics, and the engineering overhead of building bespoke reporting pipelines.

Governance is inherited from Unity Catalog automatically — there is no separate permission layer to maintain.

[→ Databricks docs](#)

PLATFORM UPDATE

Genie Chat on Workspace Home

The AI assistant previously known under the Databricks One branding has been unified and **rebranded to Genie** — the AI-powered Lakehouse assistant — now embedded on the workspace home page from first login, with no setup required.

DATA BRICKS DIGEST

MAY 2026
Lakehouse Platform Intelligence

BUSINESS VALUE

AI-assisted data exploration available from first login — no setup overhead. Genie Forge compresses the effort to scale Genie Spaces across multiple domains and business units from weeks to hours, making broad deployment of governed natural language data access operationally viable for the first time.

[→ Databricks docs](#)

LAKEFLOW CONNECT · NEW INTEGRATION

Outlook Email Ingestion

A new **managed Outlook connector in Lakeflow Connect** enables structured ingestion of Outlook email data into Unity Catalog. Supports full history or a configurable start date, incremental runs, and attachment handling — no custom connector required.

[→ Databricks docs](#)